

山东大学计算机科学与技术学院

大数据分析实践课程实验报告

学号: 202300130220	姓名: 刘傲宇	班级: 数据科学与大数据技术班
实验题目: spark 实践		
实验学时: 2	实验日期: 2025/11/14	

实验目标:

本实验旨在介绍学习者如何配置和运行 Apache Spark, 以及如何使用 Spark 进行简单的数据处理和分析。实验将涵盖以下内容:

1. 安装和配置 Apache Spark。
2. 运行一个简单的 Spark 应用程序, 以理解 Spark 的基本概念。
3. 使用 Spark 进行数据处理和分析。

实验步骤与内容:

1、环境配置

设置 spark、java 环境变量

```
hadoop@ubuntu:~$ java -version
java version "1.8.0_321"
Java(TM) SE Runtime Environment (build 1.8.0_321-b07)
Java HotSpot(TM) 64-Bit Server VM (build 25.321-b07, mixed mode)
hadoop@ubuntu:~$ export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
hadoop@ubuntu:~$ export PATH=$JAVA_HOME/bin:$PATH
hadoop@ubuntu:~$ source ~/.bashrc
hadoop@ubuntu:~$ █
```

配置、启动 spark 集群

```
hadoop@ubuntu:/usr/local/hadoop/etc/hadoop$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [ubuntu]
2025-04-16 17:10:40,142 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
hadoop@ubuntu:/usr/local/hadoop/etc/hadoop$ jps
4369 Jps
2769 org.eclipse.equinox.launcher_1.5.700.v20200207-2156.jar
3874 NameNode
4196 SecondaryNameNode
4011 DataNode
```

2、运行样例代码。

使用 PySpark 对一个销售数据集进行处理, 筛选出销售类别为 “Accessories” 的数据, 并按日期汇总每一天的销售收入。

```

1  from pyspark.sql import SparkSession
2
3  # 初始化 SparkSession
4  spark = SparkSession.builder.appName("SimpleSparkApp").getOrCreate()
5
6  # 加载数据
7  data = spark.read.csv('sales_data.csv', header=True, inferSchema=True)
8
9  # 执行一些数据处理操作
10 result = data.filter(data["Product_Category"] == "Accessories").groupBy("Date").sum("Revenue")
11
12 # 显示结果
13 result.show()

```

运行结果：

```

hadoop@ubuntu:~/Desktop/lay_temp$ python3 1.py
2025-11-17 04:07:23,054 WARN util.Utils: Your hostname, ubuntu resolves to a loopback address: 127.0.1.1; using 192.168.126.129 instead (on interface ens3)
2025-11-17 04:07:23,056 WARN util.Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
2025-11-17 04:07:24,548 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
+-----+-----+
| Date|sum(Revenue)|
+-----+-----+
| 2016/5/12|      21505|
| 2015/7/29|       5994|
| 2013/8/2|     17700|
| 2013/11/10|     20464|
| 2015/11/24|     28005|
| 2015/12/23|     19262|
| 2014/2/28|     28865|
| 2016/3/22|     16684|
| 2015/12/19|     30493|
| 2013/11/3|     26902|
| 2016/4/17|     22595|
| 2016/4/25|     19838|
| 2014/1/4|     14332|
| 2014/2/24|     19069|
| 2016/2/4|     17017|
| 2015/10/5|     25070|
| 2014/3/29|     16641|
| 2015/10/8|     22254|
| 2016/6/20|     27200|
| 2013/8/28|     25919|
+-----+-----+
only showing top 20 rows

```

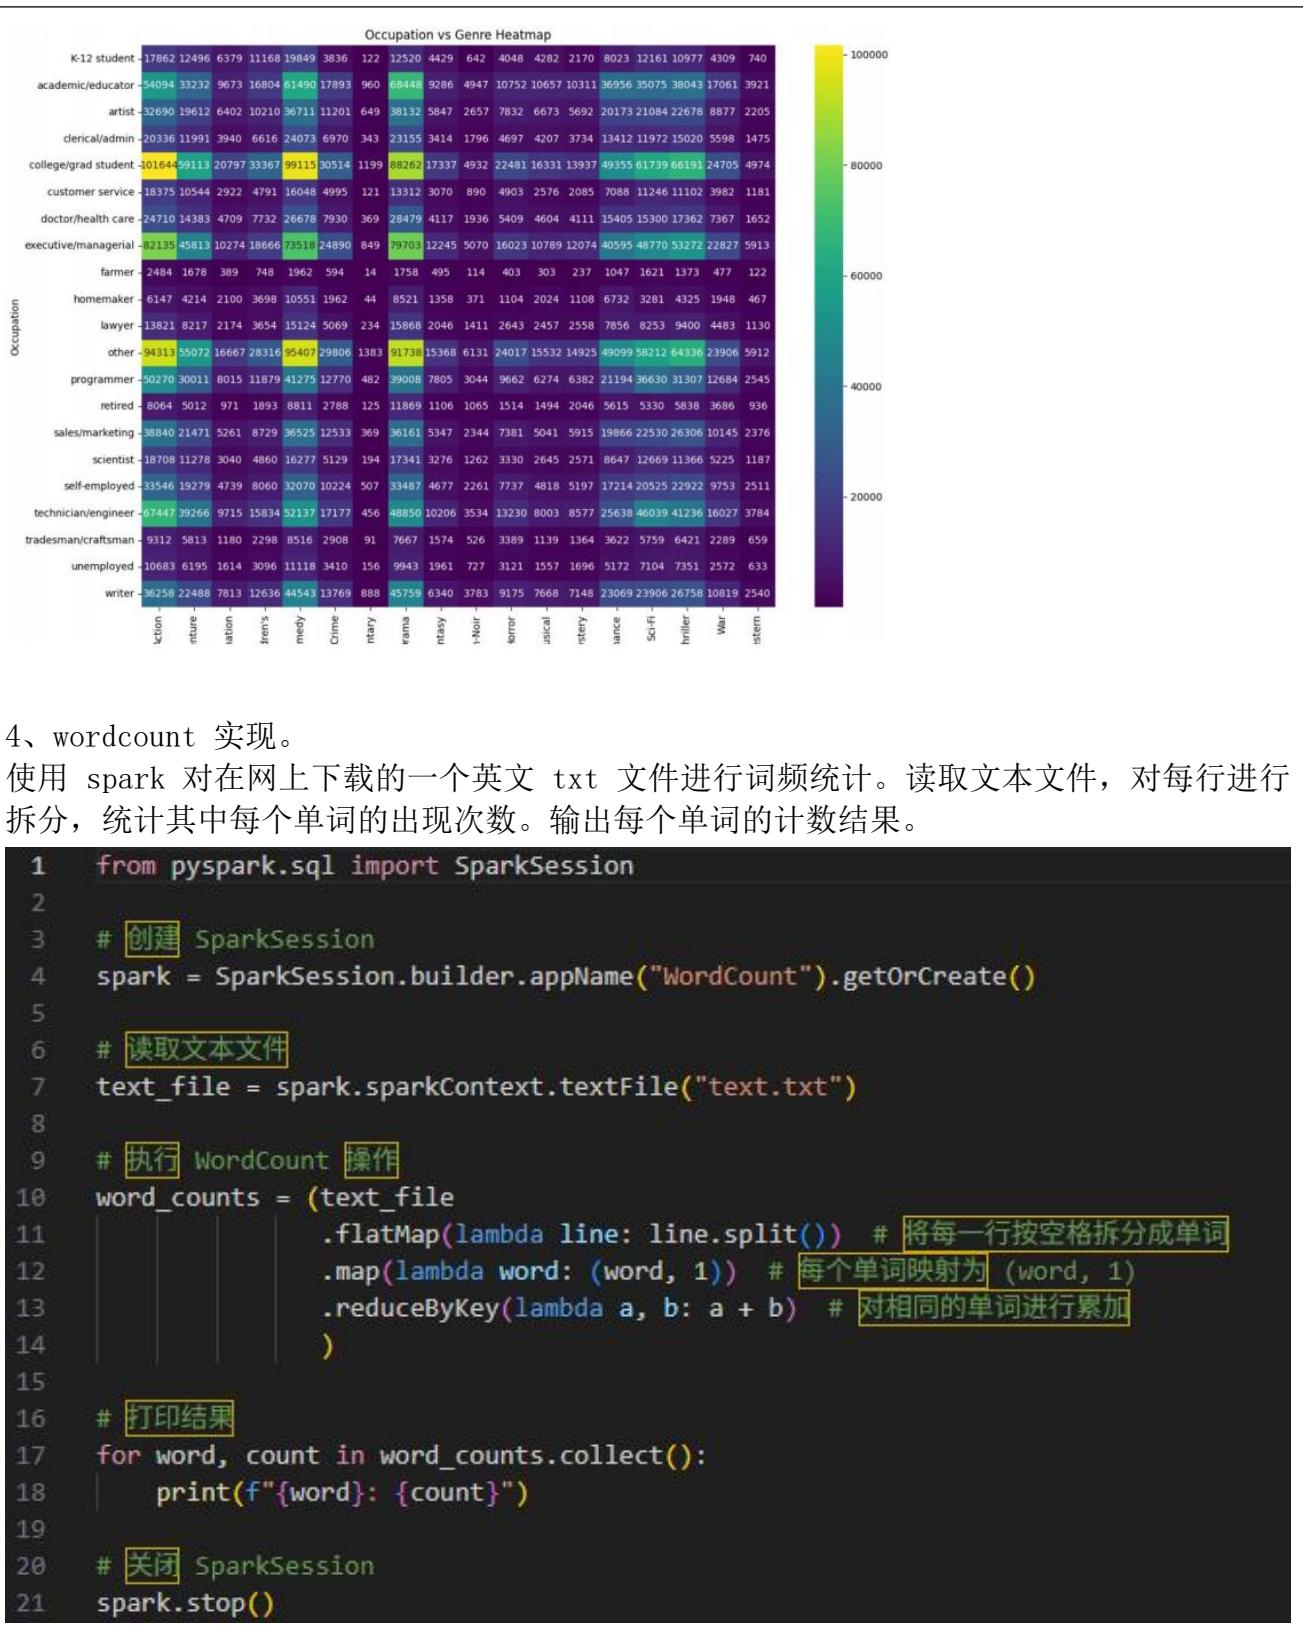
3、使用 spark 进行数据处理和分析。

分析各职业对不同种类的电影关注度，绘制热力图。

```

1  # 解析电影种类
2  movies = movies.withColumn('genre', explode(split(col('genres'), '[]')))
3
4  # 将数据连接起来
5  ratings_with_genres = ratings.join(movies, on='movieId').join(users, on='userId')
6
7  # 各类职业对不同种类电影的关注度
8  occupation_genre_counts = ratings_with_genres.groupBy('occupation', 'genre').count().toPandas()
9  occupation_genre_counts['occupation'] = occupation_genre_counts['occupation'].map(occupation_map)
10
11 # 生成透视表
12 heatmap_data = occupation_genre_counts.pivot_table(index='occupation', columns='genre', values='count', fill_value=0)
13
14 # 绘制热力图
15 plt.figure(figsize=(14, 10))
16 sns.heatmap(heatmap_data, annot=True, fmt='g', cmap='viridis')
17 plt.title('Occupation vs Genre Heatmap')
18 plt.xlabel('Genre')
19 plt.ylabel('Occupation')
20 plt.xticks(rotation=90)
21 plt.savefig('occupation_genre_heatmap.png')
22 plt.show()

```



```

hadoop@ubuntu:~/Desktop/lay_temp$ python3 2.py
2025-11-17 04:09:19,465 WARN util.Utils: Your hostname, ubuntu resolves to a loopback address: 127.0.1.1; using 192.168.126.129 instead (on interface ens3)
2025-11-17 04:09:19,468 WARN util.Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
2025-11-17 04:09:21,068 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Saying: 1
Good-bye: 1
Cambridge: 3
Again: 1
Very: 2
quietly: 4
take: 2
leave: 2
as: 2
came: 2
in: 4
The: 2
willows: 1
riverside: 1
young: 1
sun: 1
reflections: 1
shimmering: 1
waves: 2
depth: 1
of: 5
sludge: 1
Sways: 1
leisurely: 1
water: 1
gentle: 1
would: 1
Water: 2
That: 1
shade: 1
elm: 1
Holds: 1
but: 1
rainbow: 1

```

5、使用 PySpark 对销售数据进行分析 求工作日订单量最大的地区：

```

1  from pyspark.sql import SparkSession
2  from pyspark.sql.functions import to_date, dayofweek
3
4  spark = SparkSession.builder \
5      .appName("FindMaxWeekdayStoreTraffic") \
6      .config("spark.sql.legacy.timeParserPolicy", "LEGACY") \
7      .getOrCreate()
8
9  file_path = "sales_data.csv"
10 traffic_data = spark.read.csv(
11     file_path,
12     header=True, # 读取表头
13     inferSchema=True # 自动推断字段类型
14 )
15
16 max_traffic_store = (
17     traffic_data
18     .filter(dayofweek(to_date("Date", "yyyy/MM/dd")).between(2, 6)) # 筛选工作日（周一到周五：2-6）
19     .filter("Order_Quantity > 0") # 排除无效订单
20     .groupBy("State") # 按地区分组
21     .sum("Order_Quantity") # 计算总订单量
22     .withColumnRenamed("sum(Order_Quantity)", "total_order_quantity")
23     .orderBy("total_order_quantity", ascending=False) # 按总订单量降序排列
24 )
25
26 print("工作日总订单量最大的地区TOP20 (地区 总订单量): ")
27 max_traffic_store.show(20) # 显示前20名
28
29 spark.stop() # 停止Spark会话

```

```

hadoop@ubuntu:~/Desktop/lay_temp$ python3 3.py
2025-11-17 04:23:09,563 WARN util.Utils: Your hostname, ubuntu resolves to a loopback address: 127.0.1.1; using 192.168.126.129 instead (on interface ens3)
2025-11-17 04:23:09,565 WARN util.Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
2025-11-17 04:23:11,010 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
工作日总订单量最大的地区TOP20 (地区-总订单量) :
+-----+-----+
| State|total_order_quantity|
+-----+-----+
| California| 191174|
| British Columbia| 137485|
| England| 112637|
| Washington| 106059|
| New South Wales| 81750|
| Oregon| 47347|
| Victoria| 46523|
| Queensland| 41560|
| Saarland| 22387|
| Nordrhein-Westfalen| 20190|
| Seine (Paris)| 19656|
| Hessen| 18877|
| Hamburg| 14403|
| Seine Saint Denis| 14265|
| Nord| 13728|
| South Australia| 11211|
| Bayern| 10811|
| Hauts de Seine| 9165|
| Essonne| 8751|
| Yveline| 7746|
+-----+-----+
only showing top 20 rows

```

求每个地区的日均订单量：

```

1  from pyspark.sql import SparkSession
2  from pyspark.sql.functions import to_date, dayofweek
3
4  spark = SparkSession.builder \
5      .appName("CalculateStoreDailyAvgTraffic") \
6      .config("spark.sql.legacy.timeParserPolicy", "LEGACY") \
7      .getOrCreate()
8
9  file_path = "sales_data.csv"
10 traffic_data = spark.read.csv(
11     file_path,
12     header=True,
13     inferSchema=True
14 )
15
16 store_daily_avg = (
17     traffic_data
18     # 筛选工作日 (周一到周五)
19     .filter(dayofweek(to_date("Date", "yyyy/MM/dd")).between(2, 6))
20     .filter("Order_Quantity > 0") # 排除无效订单
21     .groupBy("State") # 按地区分组
22     .avg("Order_Quantity") # 计算日均订单量
23     .withColumnRenamed("avg(Order_Quantity)", "avg_order_quantity")
24     .orderBy("avg_order_quantity", ascending=False) # 按日均订单量降序排列
25 )
26
27 print("各地区工作日日均订单量TOP20 (地区-日均订单量):")
28 store_daily_avg.show(20) # 显示前20名
29
30 spark.stop() # 停止Spark会话
31

```

```

hadoop@ubuntu:~/Desktop/lay_temp$ python3 4.py
2025-11-17 04:24:31,683 WARN util.Utils: Your hostname, ubuntu resolves to a loopback address: 127.0.1.1; using 192.168.126.129 instead (on interface ens3)
2025-11-17 04:24:31,686 WARN util.Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
2025-11-17 04:24:33,192 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
各地区工作日日均订单量top20 (地区-日均订单量) :
+-----+-----+
| State|avg_order_quantity|
+-----+-----+
| Alabama|      25.5|
| Mississippi|    22.75|
| Minnesota| 22.3333333333332|
| Ohio| 20.8333333333332|
| Ontario| 20.3333333333332|
| Montana| 20.3333333333332|
| Virginia|     16.5|
| North Carolina| 16.5|
| Pas de Calais| 14.061538461538461|
| New York|      14.0|
| British Columbia| 13.666500994035784|
| Georgia| 13.571428571428571|
| Oregon| 12.734534696073158|
| Val de Marne| 12.63063063063063|
| Val d'Oise| 12.570621468926554|
| Washington| 12.304353172651254|
| Somme| 12.146067415730338|
| Essonne| 12.070344827586206|
| California| 11.973819366153075|
| Seine Saint Denis| 11.90734575959933|
+-----+-----+
only showing top 20 rows

```

结论分析与体会：

1、本实验主要依托 PySpark 框架对销售数据开展基础的数据处理与分析工作，涵盖数据加载、过滤、分组、聚合、排序等核心操作。实验结果表明，Spark 的分布式计算模型使其具备高效处理大规模数据的能力，且可在短时间内完成复杂的数据分析任务。

2、Spark 实现的优势：

高效的分布式计算能力：Spark 可将数据分布部署于多节点并执行并行计算，大幅提升计算效率，适用于大规模数据场景。

灵活的 API 体系：Spark 提供了功能强大的 DataFrame 与 RDD API，极大简化了数据分析流程。

3、Spark 实现的局限性：

在处理大规模数据时，资源的高效管理面临挑战，需规避性能瓶颈与数据倾斜问题。针对需执行复杂计算的大数据集，需对集群资源进行合理配置。